Served stats

Rendered 2026-10-15 17:32 UTC from stats/*.json, the ledger every harness run and harness export adds to (python3 -m harness stats renders it again). Aggregate tok/s, greedy, 256-token replies, the recipe’s own ruler; engine rounds from the server’s own counters.

Overview: by model and hardware

Aggregate tok/s as prose/code/structured/list; first token in seconds (one stream / the most streams).

model	hardware	one stream	most streams: total	first token (s)	suite	measured
Qwen3.8-Flash-Next INT4-Mixed (AutoRound)	2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE	69/105/86/110	0.08 /	2026-10-10		

Qwen3.8-Flash-Next INT4-Mixed (AutoRound) on 2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE

Target qwen38-int4mixed: Qwen3.8-Flash-Next INT4-Mixed, TensorFold native, TP=2 on two DGX Sparks. Engine: TensorFold native, CUDA; served as Qwen3.8-Flash-Next-INT4-Mixed. Latest: run 20261010-094527-quick (2026-10-10 09:45, plan quick, done).

Serving: patch_sha bc0402c1b6f31946984c0c469bb85d77eb926f4dc31ff9b8fcdb876ee4dbe8b5; image tf-qwen38-int4mixed:native-v1.0.4 sha256:159e8b61d9ec0827e4a2c76b9933b3ec3e358dc05ab0b7067; server_args --context 1048576 --no-thinking --parallel 16 --vision-urls --vision

ruler-fresh

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	68.8 (68.8)	105.3 (105.3)	86.0 (86.0)	109.5 (109.5)	0.08-0.11 s

ruler-after

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	67.2 (67.2)	103.5 (103.5)	83.8 (83.8)	108.4 (108.4)	0.08-0.12 s

suite

check	passed
drafts	12/12
concurrent	12/12
streaming	4/4
multiturn	8/8
stop	3/3
thinking	10/10
copy	2/2
max-tokens	3/3

checks

check	result
memory-idle	measured: head 20.5 GiB available; worker 24.9 GiB available
structured	pass: 18/18 obey; a bad regex -> 400
sampling	pass: 9/9 seeded replies equal on repeat and undrafted; 0 errors
turns	pass: resumed == again == cold: True; cached 5575 of 5642; ttft resumed 0.121s, cold 2.088s
urls	pass: public 200; loopback 400; http 400

history

run	plan	status	1 stream code	16 streams code	16 streams first token	suite
2026-10-10 09:45	quick	done	105.3			54/54