Served stats

Rendered 2026-10-15 19:05 UTC from stats/*.json, the ledger every harness run and harness export adds to (python3 -m harness stats renders it again). Aggregate tok/s, greedy, 256-token replies, the recipe's own ruler; engine rounds from the server's own counters.

Overview: by model and hardware

Aggregate tok/s as prose/code/structured/list; first token in seconds (one stream / the most streams).

model	hardware	one stream	most streams: total	first token (s)	suite	measured
GLM-5.3-Flash, EXL3 4 bpw	2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE	46/69/57/59	0.18 /	tools 59/60, json 30/30, gsm8k 246/250, humaneval 159/164, ifeval 118/150	2026-10-14	
GLM-5.3-Flash, MLX 4-bit	1x Apple M5 Ultra (256 GB unified), macOS	91/104/99/42	0.12 /		2026-10-10	
Qwen3.8-27B, MLX 4-bit + DFlash2	1x NVIDIA RTX 5090 (32 GB)	305 (median of 4 prompts)	0.10 /	tools 59/60, json 30/30, gsm8k 242/250, humaneval 153/164, ifeval 124/150	2026-10-14	
Qwen3.8-Flash-Next INT4-Mixed (AutoRound)	2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE	77/127/101/130	16: 302/510/391/519	0.03 / 0.44-0.55	tools 60/60, json 30/30, gsm8k 241/250, humaneval 157/164, ifeval 130/150	2026-10-15
Qwen3.8-Flash-Next, 6-bit with vision	1x Apple M5 Ultra (256 GB unified), macOS	128/174/148/175	8: 272/330/292/339	0.16 / 0.76-0.84	tools 59/60, json 30/30	2026-10-12
Qwen3.8-Flash-Next, MLX 6-bit	1x Apple M5 Ultra (256 GB unified), macOS	145/206/172/194	4: 566/816/673/854	0.09 / 1.63-3.07	gsm8k 242/250, ifeval 134/150	2026-10-11
Step-5-Preview, MLX 2-bit experts (8-bit rest) + MTP	1x Apple M5 Ultra (256 GB unified), macOS	39/51/41/48	0.23 /	tools 58/60, json 30/30, gsm8k 49/50, ifeval 41/50	2026-10-15	

GLM-5.3-Flash, EXL3 4 bpw on 2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE

Target glm53-exl3-sparks: GLM-5.3-Flash EXL3 4 bpw, TensorFold native, TP=2 on two DGX Sparks. Engine: TensorFold native, CUDA; served as GLM-5.3-Flash-EXL3. Latest: run evidence/2026-10-14 (2026-10-14 00:00, plan published evidence, published).

ruler-shipped-14fb445e-c1-sampled

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	45.1 (45.2)	64.6 (65.4)	56.5 (56.2)	57.0 (54.9)	0.18-0.25 s

ruler-shipped-14fb445e-c1

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	46.5 (45.4)	68.7 (69.4)	57.0 (56.4)	58.8 (55.3)	0.18-0.24 s

ruler-shipped-14fb445e-c16

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
16	144.0 (8.9)	198.3 (12.1)	180.1 (11.3)	208.4 (12.5)	0.91-1.19 s

ruler-shipped-14fb445e-c8

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
8	110.1 (13.6)	152.5 (18.6)	138.7 (17.5)	157.0 (18.7)	0.49-0.63 s

ruler-shipped-f95f2f70-c1-sampled

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token

```
  1  45.4 (45.4)  65.8 (67.0)  56.4 (55.9)  58.1 (55.3)  0.17-0.23 s

ruler-shipped-f95f2f70-c1
-------------------------

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  1  47.5 (46.5)  69.5 (69.9)  58.2 (56.9)  60.2 (56.8)  0.16-0.23 s

ruler-shipped-f95f2f70-c16
--------------------------

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  16  144.8 (8.9)  202.2 (12.3)  181.4 (11.4)  207.6 (12.5)  0.85-1.14 s

ruler-shipped-f95f2f70-c8
-------------------------

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  8  111.1 (13.7)  154.1 (18.7)  139.5 (17.2)  157.3 (18.8)  0.45-0.60 s

ablit/suite
-----------

  check  passed
  gsm8k  246/250
  mmlu  292/342
  humaneval  159/164
  ifeval  118/150
  tools  59/60
  json  30/30
  mgsm  221/240

suite-ab
--------

  check  passed
  gsm8k  244/250
  mmlu  292/342
  humaneval  156/164
  ifeval  118/150
  tools  59/60
  json  30/30
  mgsm  223/240

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-14 00:00  published evidence  published  68.7      1125/1236
  2026-10-13 00:00  published evidence  published  1083.6      1874/2172
  2026-10-12 00:00  published evidence  published  62.8  169.5  0.45-1.19 s

GLM-5.3-Flash, MLX 4-bit on 1x Apple M5 Ultra (256 GB unified), macOS
---------------------------------------------------------------------

Target glm53-m5: GLM-5.3-Flash MLX 4-bit, TensorFold native on one M5 Ultra Mac (one stream, greedy). Engine: TensorFold native, Metal; served as GLM-5.3-Flash-
    MLX-4bit. Latest: run evidence/2026-10-10-wiredreads (2026-10-10 00:00, plan published evidence, published).

ruler-repeat
------------
```

```
  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  1  92.4 (92.4)  103.6 (103.6)  97.0 (97.0)  94.9 (94.9)  0.11-0.15 s

ruler
-----

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  1  90.7 (90.7)  103.7 (103.7)  99.1 (99.1)  41.8 (41.8)  0.12-0.15 s

suite
-----

  check  passed
  snapshot  12/12
  drafts  12/12
  concurrent  12/12

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-10 00:00  published evidence  published  103.7      36/36
  2026-10-10 00:00  published evidence  published        1128/1213
  2026-10-10 00:00  published evidence  published        1137/1213
  2026-10-10 00:00  published evidence  published        1079/1213
  2026-10-09 00:00  published evidence  published  104.6
  2026-10-09 00:00  published evidence  published        349/403
  2026-10-09 00:00  published evidence  published        42/42

Qwen3.8-27B, MLX 4-bit + DFlash2 on 1x NVIDIA RTX 5090 (32 GB)
--------------------------------------------------------------

Target qwen38-27b-5090: Qwen3.8-27B MLX 4-bit + DFlash2, TensorFold native CUDA on one RTX 5090 (fp8 KV, 262k, one stream, vision). Engine: TensorFold native, C
    UDA; served as qwen38-27b. Latest: run evidence/2026-10-14-merge-1.0.11 (2026-10-14 00:00, plan published evidence, published).

bench_openai-64tok
------------------

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  1  -  -  -  -  0.10-0.10 s

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-14 00:00  published evidence  published
  2026-10-13 00:00  published evidence  published
  2026-10-12 00:00  published evidence  published        1229/1437
  2026-10-12 00:00  published evidence  published        1289/1437
  2026-10-12 00:00  published evidence  published
  2026-10-12 00:00  published evidence  published

Qwen3.8-Flash-Next, MLX 6-bit on 1x Apple M5 Ultra (256 GB unified), macOS
--------------------------------------------------------------------------

Target qwen38-flashnext-mlx6bit-m5: Qwen3.8-Flash-Next MLX 6-bit, TensorFold native on one M5 Ultra. Engine: TensorFold native, Metal; served as Qwen3.8-Flash-N
    ext-MLX-6bit. Latest: run evidence/2026-10-11-suite-rest (2026-10-11 00:00, plan published evidence, published).

suite
```

```
-----

  check  passed
  ifeval  134/150
  gsm8k  242/250
  mgsm  225/240
  mmlu  282/342
  repetition  8/8
  long-context  55/55
  concurrent  12/12
  snapshot  12/12

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-11 00:00  published evidence  published        970/1069
  2026-10-11 00:00  published evidence  published  205.9      24/24
  2026-10-11 00:00  published evidence  published        157/164
  2026-10-11 00:00  published evidence  published  199.8
  2026-10-11 00:00  published evidence  published        131/132
  2026-10-11 00:00  published evidence  published  206.2

Qwen3.8-Flash-Next, 6-bit with vision on 1x Apple M5 Ultra (256 GB unified), macOS
--------------------------------------------------------------------------------

Target qwen38-flashnext-vision-m5: Qwen3.8-Flash-Next 6-bit with image input, TensorFold native on one M5 Ultra. Engine: TensorFold native, Metal; served as Qwe
    n3.8-Flash-Next-6bit. Latest: run evidence/2026-10-12-parallel8 (2026-10-12 00:00, plan published evidence, published).

ruler
-----

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
  1  127.7 (127.7)  174.1 (174.1)  147.8 (147.8)  175.4 (175.4)  0.16-0.17 s
  4  210.2 (52.9)  303.3 (76.3)  249.2 (61.1)  316.9 (77.8)  0.37-0.40 s
  8  271.5 (34.1)  330.0 (40.8)  292.5 (36.5)  339.3 (41.7)  0.76-0.84 s

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-12 00:00  published evidence  published  174.1
  2026-10-12 00:00  published evidence  published  202.7
  2026-10-12 00:00  published evidence  published  176.6
  2026-10-11 00:00  published evidence  published  205.6      131/132
  2026-10-11 00:00  published evidence  published  178.9      513/544

Qwen3.8-Flash-Next INT4-Mixed (AutoRound) on 2x NVIDIA DGX Spark (GB10, 128 GB unified each), TP=2 over RoCE
-----------------------------------------------------------------------------------------------------------

Target qwen38-int4mixed: Qwen3.8-Flash-Next INT4-Mixed, TensorFold native, TP=2 on two DGX Sparks. Engine: TensorFold native, CUDA; served as Qwen3.8-Flash-Next
    -INT4-Mixed. Latest: run 20261015-171328-full (2026-10-15 17:13, plan full, done).

Serving: patch_sha 186193c6a530d89ee6383c11a9e41f98bac95bb119959aa5901c65b0f93b21ea; image tf-qwen38-int4mixed:native-v1.0.11-3 sha256:7c5838987336819a3dcc51d8f
    9abcf169f80fbfe73222e; server_args --context 1048576 --no-thinking --parallel 16 --vision-urls --vision

ruler-fresh
-----------

  streams  prose total (a stream)  code total (a stream)  structured total (a stream)  list total (a stream)  first token
```

1	77.3 (77.3)	127.0 (127.0)	100.6 (100.6)	129.6 (129.6)	0.03-0.04 s
4	159.1 (39.8)	253.9 (65.0)	201.7 (49.8)	268.4 (61.8)	0.07-0.10 s
8	228.9 (28.8)	385.2 (46.2)	282.4 (35.8)	423.1 (51.3)	0.11-0.14 s
16	301.9 (18.8)	509.8 (30.6)	391.3 (24.0)	519.2 (31.6)	0.44-0.55 s

Engine rounds (ruler-fresh):

streams	tokens a round	drafts a round	accepted	ms a round	requests
1	3.621	4.042	66%	33.33	12
4	3.173	2.949	74%	58.64	49
8	3.09	2.653	80%	76.26	100
16	3.222	2.925	77%	122.65	196

ruler-after

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	77.5 (77.5)	125.9 (125.9)	99.2 (99.2)	129.7 (129.7)	0.04-0.04 s
16	302.3 (18.8)	509.5 (30.4)	389.8 (23.8)	520.4 (31.6)	0.44-0.50 s

Engine rounds (ruler-after):

streams	tokens a round	drafts a round	accepted	ms a round	requests
1	3.567	4.016	65%	33.01	13
16	3.21	2.914	76%	122.7	193

suite

check	passed
drafts	12/12
concurrent	12/12
streaming	4/4
multiturn	8/8
stop	3/3
thinking	10/10
copy	2/2
max-tokens	3/3
tools	60/60
json	30/30
ifeval	130/150
gsm8k	241/250
mgsm	219/240
mmlu	283/342
humaneval	157/164
repetition	8/8
long-context	55/55
vision	12/12
soak	60/60
snapshot	12/12

checks

check	result
memory-idle	measured: head 21.8 GiB available; worker 25.2 GiB available
structured	pass: 18/18 obey; a bad regex -> 400
sampling	pass: 9/9 seeded replies equal on repeat and undrafted; 0 errors
turns	pass: resumed == again == cold: True; cached 5575 of 5642; ttft resumed 0.118s, cold 0.035s
urls	pass: public 200; loopback 400; http 400

capacity pass: 4/6 answered, 2 refused, 0 lost (no reply); a short request after: 200

history
----

run	plan	status	1 stream code	16 streams code	16 streams first token	suite
2026-10-15 17:13	full	done	127.0	509.8	0.44-0.55 s	1321/1437
2026-10-15 00:00	published evidence	published	123.4	472.9	0.52-0.74 s	1321/1437
2026-10-15 00:00	published evidence	published		1321/1437		
2026-10-11 00:00	published evidence	published	120.1	507.2	0.54-0.62 s	1321/1437
2026-10-11 00:00	published evidence	published	120.4	507.1	0.50-0.59 s	
2026-10-10 09:45	quick	done	105.3	54/54		
2026-10-10 00:00	published evidence	published	122.0	510.3	0.40-0.49 s	114/114
2026-10-10 00:00	published evidence	published	125.9	510.9	0.44-0.49 s	1261/1377
2026-10-10 00:00	published evidence	published	125.7	506.7	0.42-0.49 s	54/54
2026-10-10 00:00	published evidence	published	106.1	506.1	0.42-0.51 s	54/54
2026-10-10 00:00	published evidence	published	110.0	511.8	0.41-0.55 s	1261/1377
2026-10-10 00:00	published evidence	published	105.8	505.5	0.41-0.46 s	1261/1377
2026-10-09 00:00	published evidence	published	105.2	412.2	0.78-0.98 s	1260/1377
2026-10-09 00:00	published evidence	published		1256/1377		
2026-10-08 00:00	published evidence	published	107.1	519.4	0.44-0.56 s	
2026-10-08 00:00	published evidence	published	107.7	474.0	0.52-0.66 s	145/145
2026-10-08 00:00	published evidence	published	107.6	442.6	0.44-0.58 s	36/36
2026-10-08 00:00	published evidence	published	107.6	441.7	0.45-0.58 s	
2026-10-07 00:00	published evidence	published		12/12		

Step-5-Preview, MLX 2-bit experts (8-bit rest) + MTP on 1x Apple M5 Ultra (256 GB unified), macOS

Target step5-m5: Step-5-Preview 2-bit (experts 2/128, rest 8/64, MTP, images), TensorFold native Metal on one M5 Ultra. Engine: TensorFold native, Metal; served as step5-preview. Latest: run 20261015-134330-one-stream (2026-10-15 13:43, plan one-stream, done).

Serving: engine_commit aa221b11; binary_sha256 4f89cf81182d203e84aac48335e84a90a29ba6ae26de96b32ec387e50c7e878f; served_model {"object": "list", "data": [{"id": "step5-preview", "object": "model", "created": 17920898; chip Apple M5 Ultra
274877906944

ruler-fresh
----

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	39.2 (39.2)	50.6 (50.6)	40.7 (40.7)	48.2 (48.2)	0.23-0.27 s

Engine rounds (ruler-fresh):

streams	tokens a round	drafts a round	accepted	ms a round	requests
1	2.381	2.072	67%	53.3	12

ruler-after
----

streams	prose total (a stream)	code total (a stream)	structured total (a stream)	list total (a stream)	first token
1	39.2 (39.2)	50.6 (50.6)	40.6 (40.6)	48.1 (48.1)	0.23-0.27 s

Engine rounds (ruler-after):

streams	tokens a round	drafts a round	accepted	ms a round	requests
1	2.339	1.939	69%	52.42	13

suite

```
  check  passed
  drafts  12/12
  streaming  4/4
  multiturn  8/8
  stop  3/3
  thinking  10/10
  copy  2/2
  max-tokens  3/3
  tools  58/60
  json  30/30
  gsm8k  49/50
  ifeval  41/50

checks
------

  check  result
  memory-idle  measured: mac: System-wide memory free percentage: 95%
total = 2048.00M  us |
  sampling  pass: 9/9 seeded replies equal on repeat and undrafted; 0 errors
  memory-busy  measured: mac: System-wide memory free percentage: 28%
total = 2048.00M  us |

history
-------

  run  plan  status  1 stream code  16 streams code  16 streams first token  suite
  2026-10-15 13:43  one-stream  done  50.6      220/232

Served earlier, numbers in their READMEs only
---------------------------------------------

These recipes predate the ruler's JSON, so their measurements are not in the ledger; their READMEs hold them.

  model  hardware  engine  recipe
  Qwen3.8-Flash-Next, turboderp EXL3 4.05 bpw  2x NVIDIA DGX Spark (GB10), TP=2  TensorFold 0.6.6 (Python) with an EXL3 patch  github.com/BobClawblaw/qwen38-fla
    shnext-exl3-2x-dgx-sparks
  Qwen3.8-Flash-Next (TensorFold 0.6.6 recipe)  2x NVIDIA DGX Spark (GB10), TP=2  TensorFold 0.6.6 (Python)  github.com/BobClawblaw/qwen38-flashnext-tensorfold-
    2x-dgx-sparks

Step-5-Preview on one M5 Ultra: measured outside the harness
------------------------------------------------------------

TensorFold native Metal, branch step5-native (engine 5ad7660c, images aa221b11), MLX 2-bit experts with the rest
8-bit, greedy, one stream. The harness's step5-m5 runs hold the ruler and the suite; these are the rest.

  what  measured
  prompt, 4k / 8k / 16k / 32k tokens  690 / 668 / 647 / 622 tok/s (the sparse indexer caps each row's keys past 4k)
  131,016-token chat question about its first page  answered exactly, 264 s end to end, at least 23% of memory free
  decode without drafts, 20 / 600 / 4,200 / 32,768 tokens of context  29.6 / 21.9 / 14.7 / 14.5 tok/s
  MTP drafts against '"draft": false' (replies token-identical)  code 25.7 to 44.6 tok/s (70% accepted), prose 26.4 to 34.8 (56%), math 26.2 to 46.0 (75%)
  an image, 1200 x 800 PNG (six 504-px crops and the overview, 692 prompt tokens)  2.1 s to a 4-token answer
  an image, 600 x 400 PNG (overview only, 193 prompt tokens)  0.7 s
```